



Unsupervised Algorithm - Clustering

Afizan Azman

<https://youtu.be/4b5d3muPQmA>

Contents

- Unsupervised learning
- Clustering
- Distance/similarity measures



Types of Machine Learning

Supervised

data points have known outcome

Unsupervised

data points have unknown outcome



Types of Machine Learning

Supervised

data points have known outcome

Unsupervised

data points have unknown outcome



Types of Unsupervised Learning

Clustering


identify unknown structure in data

Dimensionality
Reduction

use structural characteristics to simplify data




Unsupervised Learning Overview






Clustering: Finding Distinct Groups






Dimensionality Reduction: Simplifying Structure



Cluster

- The process of partitioning a set of data into a set of meaningful (hopefully) sub-classes, called *clusters*
 - collection of data points that are “similar” to one another and collectively should be treated as group
 - as a collection, are sufficiently different from other groups

Cluster



Clustering in






Introduction to Unsupervised Learning

Users of a web application:

One feature (age)



Age




Introduction to Unsupervised Learning

Users of a web application:

One feature (age)

Two clusters



Age




Introduction to Unsupervised Learning

Users of a web application:

One feature (age)

Three clusters



Age




Introduction to Unsupervised Learning

Users of a web application:

One feature (age)

Five clusters




Age



K-Means Algorithm


$K = 2$ (find two clusters)





K-Means Algorithm


$K = 2$, Randomly assign cluster centers





K-Means Algorithm


$K = 2$, Each point belongs to closest center





K-Means Algorithm


$K = 2$, Move each center to cluster's mean





K-Means Algorithm


$K = 2$, Each point belongs to closest center





K-Means Algorithm


$K = 2$, Move each center to cluster's mean





K-Means Algorithm


$K = 2$, Points don't change $\rightarrow$ Converged





K-Means Algorithm


$K = 2$, Each point belongs to closest center





K-Means Algorithm


$K = 3$






K-Means Algorithm

$K = 3$, Results depend on initial cluster assignment





Which Model is the Right One?





Which Model is the Right One?

- **Inertia:** sum of squared distance from each point (x_i) to its cluster (C_k)


$$\sum_{i=1}^n (x_i - C_k)^2$$

- Smaller value corresponds to tighter clusters
- Other metrics can also be used



Which Model is the Right One?


Initiate multiple times,
take model with the best score





Which Model is the Right One?


Inertia = 12.645





Which Model is the Right One?


Inertia = 12.943






Which Model is the Right One?

Inertia = 13.112






Smarter Initialization of K-Means Clusters





Smarter Initialization of K-Means Clusters


Pick one point at random as initial point





Smarter Initialization of K-Means Clusters


Pick next point with $1/\text{distance}^2$ probability





Smarter Initialization of K-Means Clusters


Pick next point with $1/\text{distance}^2$ probability





Smarter Initialization of K-Means Clusters


Pick next point with $1/\text{distance}^2$ probability





Smarter Initialization of K-Means Clusters

Assign clusters





Choosing the Right Number of Clusters



Choosing the Right Number of Clusters

- Sometimes the question has a K



Choosing the Right Number of Clusters

- Sometimes the question has a K



Choosing the Right Number of Clusters

- Sometimes the question has a K
- Clustering similar jobs on 4 CPU cores ($K=4$)



Choosing the Right Number of Clusters

- Sometimes the question has a K
- Clustering similar jobs on 4 CPU cores ($K=4$)
- A clothing design in 10 different sizes to cover most people ($K=10$)




Choosing the Right Number of Clusters

- Sometimes the question has a K
- Clustering similar jobs on 4 CPU cores ($K=4$)
- A clothing design in 10 different sizes to cover most people ($K=10$)
- A navigation interface for browsing scientific papers with 20 disciplines ($K=20$)



Choosing the Right Number of Clusters




- Inertia measures distance of point to cluster

r



Choosing the Right Number of Clusters



- Inertia measures distance of point to cluster
- Value decreases with increasing K as long as cluster density increases



K-Means: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import KMeans
```



K-Means: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import KMeans
```

Create an instance of the class

```
kmeans = KMeans(n_clusters=3,  
                 init='k-means++')
```




K-Means: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import KMeans
```

Create an instance of the class

```
kmeans = KMeans(n_clusters=3,  
                 init='k-means++')
```



final number
of clusters




K-Means: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import KMeans
```

Create an instance of the class

```
kmeans = KMeans(n_clusters=3,  
                 init='k-means++')
```



kmeans++
cluster
initiation



K-Means: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import KMeans
```

Create an instance of the class

```
kmeans = KMeans(n_clusters=3,  
                 init='k-means++')
```

Fit the instance on the data and then predict clusters for new data

```
kmeans = kmeans.fit(X1)
```



K-Means: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import KMeans
```

Create an instance of the class

```
kmeans = KMeans(n_clusters=3,  
                 init='k-means++')
```

Fit the instance on the data and then predict clusters for new data

```
kmeans = kmeans.fit(X1)  
y_predict = kmeans.predict(X2)
```

Can also be used in batch mode with [MiniBatchKMeans](#).

Basic Questions

- Data preparation - getting/setting up data for clustering
 - extraction
 - normalization
- Similarity/Distance measure - how is the distance between points defined
- Use of domain knowledge (prior knowledge)
 - can influence preparation, Similarity/Distance measure
- Efficiency - how to construct clusters in a reasonable amount of time

Similarity (Distance) Thresholds

- A similarity (distance) threshold may be used to mark pairs that are “sufficiently” similar

	T1	T2	T3	T4	T5	T6	T7
T2	7						
T3	16	8					
T4	15	12	18				
T5	14	3	6	6			
T6	14	18	16	18	6		
T7	9	6	0	6	9	2	
T8	7	17	8	9	3	16	3

	T1	T2	T3	T4	T5	T6	T7
T2	0						
T3	1	0					
T4	1	1	1				
T5	1	0	0	0			
T6	1	1	1	1	0		
T7	0	0	0	0	0	0	
T8	0	1	0	0	0	1	0


Using a threshold value of 10 in the previous example



Distance Metrics




Distance Metric Choice



- Choice of distance metric is extremely important to clustering success
- Each metric has strengths and most appropriate use-cases...
- ...but sometimes choice of distance metric is also based on empirical evaluation




Euclidean Distance






Euclidean Distance






Euclidean Distance (L2 Distance)






Manhattan Distance (L1 or City Block Distance)






Cosine Distance





Cosine Distance



Euclidean vs Cosine Distance

- Euclidean is useful for coordinate based measurements



Euclidean vs Cosine Distance

- Euclidean is useful for coordinate based measurements
- Cosine is better for data such as text where location of occurrence is less important



Euclidean vs Cosine Distance

- Euclidean is useful for coordinate based measurements
- Cosine is better for data such as text where location of occurrence is less important
- Euclidean distance is more sensitive to curse of dimensionality (see lesson 12)



Jaccard Distance

Applies to sets (like word occurrence)

- **Sentence A:** “I like chocolate ice cream.”
- set A = {I, like, chocolate, ice, cream}
- **Sentence B:** “Do I want chocolate cream or vanilla cream?”
- set B = {Do, I, want, chocolate, cream, or, vanilla}

$$1 - \frac{A \cap B}{A \cup B} = 1 - \frac{\text{len}(\text{shared})}{\text{len}(\text{unique})}$$



Jaccard Distance

Applies to sets (like word occurrence)

- **Sentence A:** “I like chocolate ice cream.”
- set A = {**I**, like, **chocolate**, ice, **cream**}
- **Sentence B:** “Do I want chocolate cream or vanilla cream?”
- set B = {Do, **I**, want, **chocolate**, **cream**, or, vanilla}

$$1 - \frac{A \cap B}{A \cup B} = 1 - \frac{3}{9}$$



Distance Metrics: The Syntax

Import the general pairwise distance function

```
from sklearn.metrics import pairwise_distances
```



Distance Metrics: The Syntax

Import the general pairwise distance function

```
from sklearn.metrics import pairwise_distances
```

Calculate the distances

```
dist = pairwise_distances(X, Y,  
                           metric='euclidean')
```




Distance Metrics: The Syntax

Import the general pairwise distance function

```
from sklearn.metrics import pairwise_distances
```

Calculate the distances

```
dist = pairwise_distances(X, Y,  
                           metric='euclidean')
```



distance metric
choice



Distance Metrics: The Syntax

Import the general pairwise distance function

```
from sklearn.metrics import pairwise_distances
```

Calculate the distances

```
dist = pairwise_distances(X, Y,  
                           metric='euclidean')
```

Other distance metric choices are: cosine, manhattan, jaccard, etc.



Distance Metrics: The Syntax

Import the general pairwise distance function

```
from sklearn.metrics import pairwise_distances
```

Calculate the distances

```
dist = pairwise_distances(X, Y,  
                           metric='euclidean')
```

Other distance metric choices are: cosine, manhattan, jaccard, etc.

Distance metric methods can also be imported specifically, e.g.:


```
from sklearn.metrics import euclidean_distances
```



Hierarchical Agglomerative Clustering




Hierarchical Agglomerative Clustering





Hierarchical Agglomerative Clustering


Find closest pair, merge into a cluster





Hierarchical Agglomerative Clustering


Find next closest pair and merge





Hierarchical Agglomerative Clustering


Find next closest pair and merge





Hierarchical Agglomerative Clustering


Keep merging closest pairs





Hierarchical Agglomerative Clustering


If the closest pair is two clusters, merge them





Hierarchical Agglomerative Clustering


Keep merging closest pairs and clusters





Hierarchical Agglomerative Clustering


Keep merging closest pairs and clusters





Hierarchical Agglomerative Clustering


Current number of clusters = 6





Hierarchical Agglomerative Clustering


Current number of clusters = 5





Hierarchical Agglomerative Clustering


Current number of clusters = 4





Hierarchical Agglomerative Clustering


Current number of clusters = 3





Hierarchical Agglomerative Clustering


Current number of clusters = 2





Hierarchical Agglomerative Clustering

Current number of clusters = 1





Agglomerative Clustering Stopping Conditions

Condition 1

the correct number of clusters is reached



Agglomerative Clustering Stopping Conditions

Condition 1

the correct number of clusters is reached


Condition 2

minimum average cluster distance
reaches a set value



Hierarchical Agglomerative Clustering


Current number of clusters = 5





Hierarchical Agglomerative Clustering


Current number of clusters = 5





Hierarchical Agglomerative Clustering


Current number of clusters = 4





Hierarchical Agglomerative Clustering


Current number of clusters = 4





Hierarchical Agglomerative Clustering


Current number of clusters = 3





Hierarchical Agglomerative Clustering


Current number of clusters = 3





Hierarchical Agglomerative Clustering


Current number of clusters = 2





Hierarchical Agglomerative Clustering


Current number of clusters = 2





Hierarchical Linkage Types


Single linkage: minimum pairwise distance between clusters





Hierarchical Linkage Types


Single linkage: minimum pairwise distance between clusters





Hierarchical Linkage Types


Complete linkage: maximum pairwise distance between clusters





Hierarchical Linkage Types


Complete linkage: maximum pairwise distance between clusters





Hierarchical Linkage Types


Average linkage: average pairwise distance between clusters





Hierarchical Linkage Types


Average linkage: average pairwise distance between clusters





Hierarchical Linkage Types


Ward linkage: merge based on best inertia





Hierarchical Linkage Types

Ward linkage: merge based on best inertia





Agglomerative Clustering: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import AgglomerativeClustering
```

Create an instance of the class

```
agg = AgglomerativeClustering(n_clusters=3,  
                               affinity='euclidean',  
                               linkage='ward')
```

Fit the instance on the data and then predict clusters for new data

```
agg = agg.fit(X1)  
y_predict = agg.predict(X2)
```




Agglomerative Clustering: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import AgglomerativeClustering
```

Create an instance of the class

```
agg = AgglomerativeClustering(n_clusters=3,  
                              affinity='euclidean',  
                              linkage='ward')
```



final number
of clusters

Fit the instance on the data and then predict clusters for new data

```
agg = agg.fit(X1)  
y_predict = agg.predict(X2)
```




Agglomerative Clustering: The Syntax

Import the class containing the clustering method

```
from sklearn.cluster import AgglomerativeClustering
```

Create an instance of the class

```
agg = AgglomerativeClustering(n_clusters=3,  
                              affinity='euclidean',  
                              linkage='ward')
```




cluster affinity
and
aggregation

Fit the instance on the data and then predict clusters for new data

```
agg = agg.fit(X1)  
y_predict = agg.predict(X2)
```

Other Types of Clustering



Reference: http://scikit-learn.org/stable/auto_examples/cluster/plot_cluster_comparison.html



Wrap Up

Thank you

QUESTIONS?